

Analyzing EEG data using GAMs

Lecture 4 of advanced regression for linguists

Martijn Wieling

Department of Humanities Computing, University of Groningen

Cambridge, January 14, 2015



This lecture

- ▶ Introduction
 - ▶ ERPs to study gender violations
 - ▶ Research question
- ▶ Design
- ▶ Methods: R code
- ▶ Results
- ▶ Discussion

ERPs to study grammatical gender violations

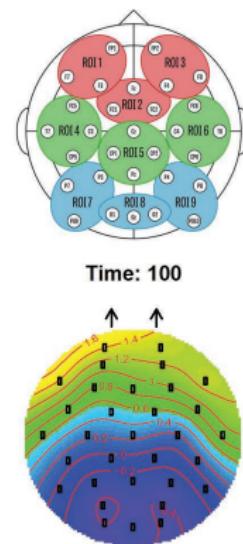
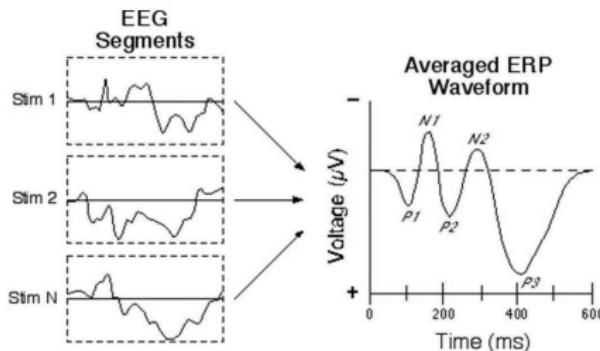
- ▶ A P600 (a positivity 'around' 600 ms. after stimulus onset) is sensitive to grammatical violations
- ▶ An N400 (a negativity 'around' 400 ms. after stimulus onset) is modulated by semantic context and lexical properties of a word
- ▶ The P600/N400 are found by **comparing** the incorrect sentences to the correct sentences
- ▶ Native speakers appear to show a P600 for gender violations
 - ▶ But analyzed by **averaging** over items and over subjects!
- ▶ Here we are interested in how non-native speakers respond to gender violations (joint work with **Nienke Meulman**)
- ▶ Gender is very hard to learn for L2 learners
- ▶ Even though behaviorally L2 learners might show correct responses, the brain may reveal differences in processing gender

Research question

- ▶ Is the P600 for gender violations dependent on age of arrival for the L2 learners of German?

ERP data

- ▶ Today: analysis of single region of interest (ROI 8)
 - ▶ GAMs allow for spatial distribution analyses



Design

- ▶ 67 L2 speakers of German (Slavic L1)
- ▶ Auditory presentation of correct sentences or sentences with a gender violation (incorrect determiner; **no** determiners in L1)
- ▶ 48 items in each condition: 96 trials per participant (minus artifacts)
- ▶ Example:

*Nach der Schlägerei ist das/*der Auge des Angestellten von der Krankenschwester versorgt worden.*

After the fight the_{neut}/*the_{masc} eye of the worker was treated by the nurse.



Data overview

```
> load('dat.rda')

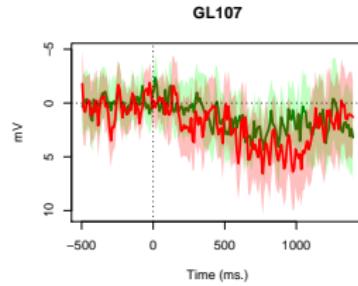
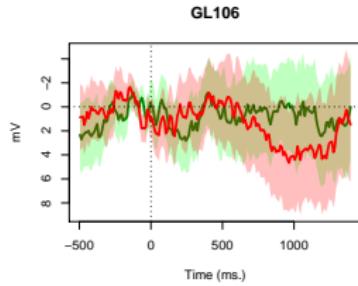
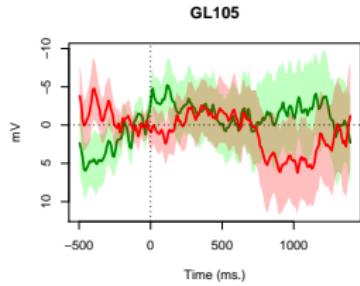
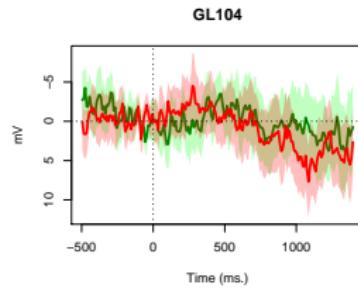
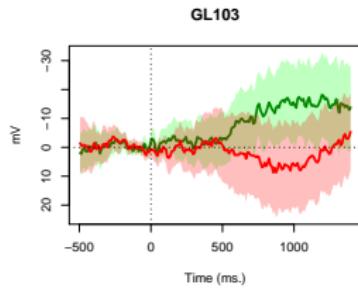
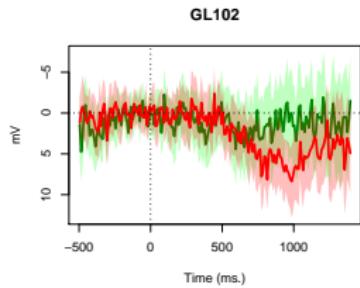
# data needs to be sorted to test for autocorrelation problems
> dat = dat[order(dat$Subject, dat$TrialNr, dat$Time, decreasing=F),]

> head(dat)
      uV Time Subject     Group Word TrialNr      Roi
44947 8.94400 505    GL102 GenEarly Wald       2 post.mid
28121 15.56267 515    GL102 GenEarly Wald       2 post.mid
40909 21.30807 525    GL102 GenEarly Wald       2 post.mid
1294   13.31573 535    GL102 GenEarly Wald       2 post.mid
9446   19.10700 545    GL102 GenEarly Wald       2 post.mid
81324 17.95607 555    GL102 GenEarly Wald       2 post.mid
          Structure Correctness L1 AoArr LoR Age Edu SeqStart
44947        DN      incor PL    8  24  32   3 TRUE
28121        DN      incor PL    8  24  32   3 FALSE
40909        DN      incor PL    8  24  32   3 FALSE
1294         DN      incor PL    8  24  32   3 FALSE
9446         DN      incor PL    8  24  32   3 FALSE
81324        DN      incor PL    8  24  32   3 FALSE

> dim(dat)
[1] 442160      15
```

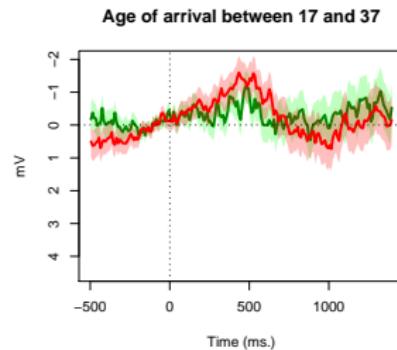
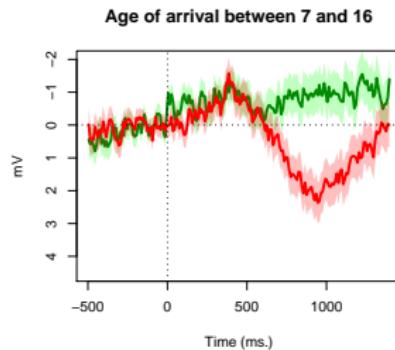
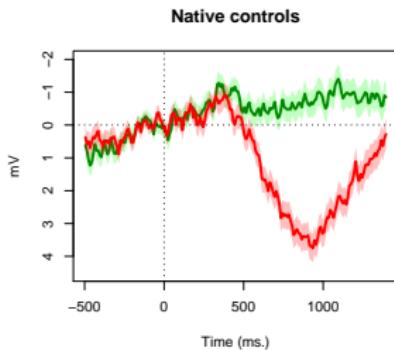
Much individual variation

(the signal has been downsampled to 100 Hz)



General patterns exist

(note the arbitrary age split, however)



A first model: a general effect of time

(we focus here on the time span from 500 to 1300 ms)

```
> library(mgcv) # 1.8.3
> m0 = bam(uV ~ s(Time), data=dat, gc.level=2, method='ML')
> summary(m0)
```

Family: gaussian

Link function: identity

Formula:

uV ~ s(Time)

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.06120	0.02774	-2.206	0.0274 *

Approximate significance of smooth terms:

edf	Ref.df	F	p-value
s(Time)	3.917	4.84	25.95 <2e-16 ***

R-sq.(adj) = 0.000284 Deviance explained = 0.0293%
ML score = 1.9162e+06 Scale est. = 340.13 n = 442160

- ▶ Why was maximum-likelihood (ML) estimation used for fitting?

A first model: a general effect of time

(we focus here on the time span from 500 to 1300 ms)

```
> library(mgcv) # 1.8.3
> m0 = bam(uV ~ s(Time), data=dat, gc.level=2, method='ML')
> summary(m0)
```

Family: gaussian

Link function: identity

Formula:

uV ~ s(Time)

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.06120	0.02774	-2.206	0.0274 *

Approximate significance of smooth terms:

edf	Ref.df	F	p-value
s(Time)	3.917	4.84	25.95 <2e-16 ***

R-sq.(adj) = 0.000284 Deviance explained = 0.0293%
ML score = 1.9162e+06 Scale est. = 340.13 n = 442160

- ▶ Why was maximum-likelihood (ML) estimation used for fitting?

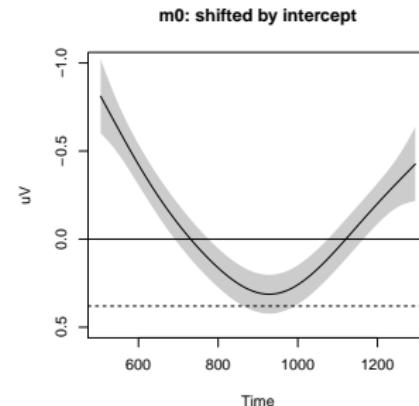
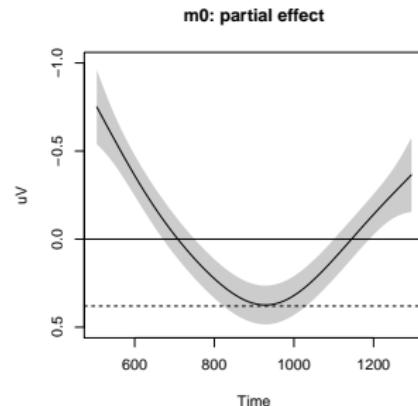
Fitting methods for GAMs

- ▶ For model comparison, the approach is similar as for `lmer` models:
 - ▶ For comparing differences in the fixed effects (constant random effects) of two models, fit both models using `method="ML"` (`lmer`: REML=F)
 - ▶ For comparing random effects (constant fixed effects) between two models, fit using `method="fREML"` (bam only) or `method="REML"` (`lmer`: REML=T)
- ▶ Choosing the fitting method of your final model (using `method="..."`):
 - ▶ ML: conservative, but biased variance components (oversmoothing)
 - ▶ fREML/REML: better smooths, but less conservative *p*-values
 - ▶ GCV.Cp: good for prediction, but not robust to autocorrelation problems
- ▶ I generally use `method="ML"` for my final model (conservative)
 - ▶ (but note that fitting a model with non-problematic residuals is currently only possible by using `method="fREML"`)

Visualizing the non-linear time pattern of the model

(Interpreting GAM results **always** involves visualization)

```
> par(mfrow=c(1,2)) # 2 plots in one window
> plot(m0, rug=F, shade=T, ylim=c(0.5,-1), seWithMean=T,
       ylab='uV', main='m0: partial effect')
> abline(h = c(0,0.38), lty = c(1,2)) # horizontal lines
> ic = m0$coef["(Intercept)"] # intercept (-0.06) from model
> plot(m0, rug=F, shade=T, ylim=c(0.5-ic,-1-ic), seWithMean=T, shift = ic,
       ylab='uV', main='m0: shifted by intercept')
> abline(h = c(0,0.38), lty = c(1,2)) # horizontal lines
```



Check for additional wigglyness

(if p-value is low and edf close to k'; note that the default k is arbitrary and higher for s than for te)

```
> gam.check(m0)
```

```
Method: ML    Optimizer: outer newton
full convergence after 2 iterations.
Gradient range [-0.6306798,0.1689593]
(score 1916153 & scale 340.1335).
Hessian positive definite, eigenvalue range [1.948326,221080.6].
Model rank = 10 / 10
```

Basis dimension (k) checking results. Low p-value (k-index<1) may indicate that k is too low, especially if edf is close to k'.

	k'	edf	k-index	p-value
s (Time)	9.000	3.917	0.997	0.46

Increasing the wigglyness of a smooth with k

(double k if higher k is needed, but do not set it too high, i.e. max $\frac{1}{2} \times$ unique time points)

```
> m0b = bam(uV ~ s(Time, k=20), data=dat, gc.level=2, method='ML')
> summary(m0b)
```

Family: gaussian

Link function: identity

Formula:

```
uV ~ s(Time, k = 20)
```

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.06119	0.02774	-2.205	0.0274 *

Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
s(Time)	4.038	5.04	24.94	<2e-16 *** # was 3.917

R-sq.(adj) = 0.000284 Deviance explained = 0.0293%
ML score = 1.9162e+06 Scale est. = 340.13 n = 442160

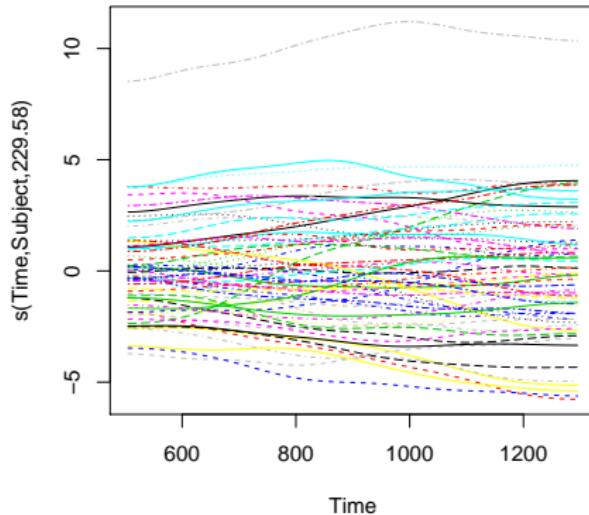
Taking into account individual variation

- ▶ The effect of time is non-linear and is variable per subject
- ▶ We need a random intercept/slope which is non-linear
 - ▶ (instead of random intercept and random slope of time per subject)

```
> m1 = bam(uV ~ s(Time) + s(Time, Subject, bs='fs', m=1),  
           data=dat, gc.level=2, method='ML')  
> summary(m1)  
...  
Parametric coefficients:  
              Estimate Std. Error t value Pr(>|t|)  
(Intercept) -0.08612    0.29531   -0.292    0.771  
  
Approximate significance of smooth terms:  
             edf Ref.df     F p-value  
s(Time)        3.741  4.482 11.49 6.09e-10 ***  
s(Time, Subject) 229.578 602.000 14.48 < 2e-16 ***  
  
R-sq.(adj) =  0.0196  Deviance explained =  2.01% # was 0.0293%  
-ML = 1.9121e+06  Scale est. = 333.56    n = 442160
```

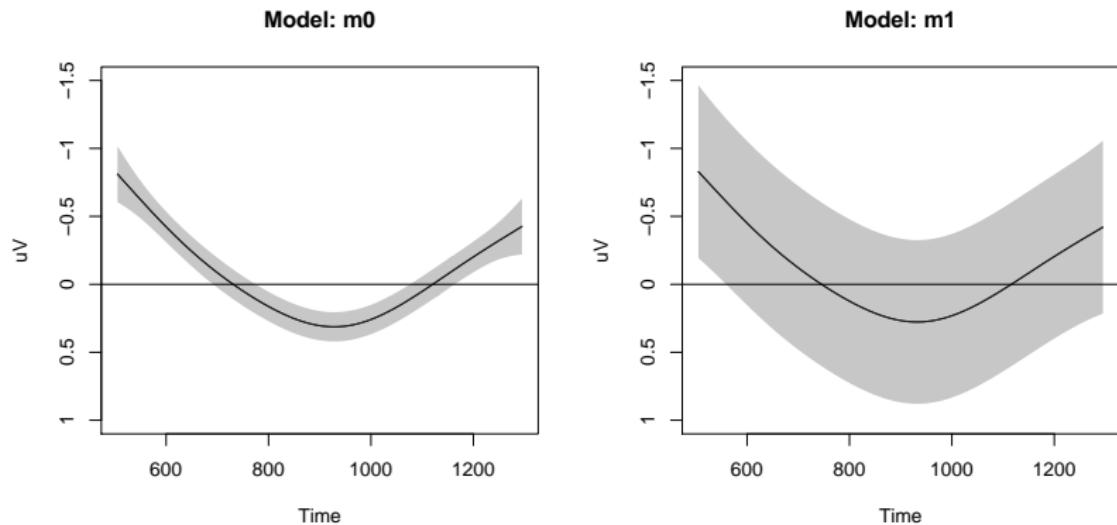
Visualization of individual variation

```
> plot(ml, select=2)
```



Smooths become more uncertain

```
> source('plotting.R') # custom plotting functions
> par(mfrow=c(1,2))
> plotSmooths(m0, "Time", ylim=c(1,-1.5)) # using complete model context
> plotSmooths(m1, "Time", ylim=c(1,-1.5), dropRanef="Subject")
```



Assessing correct versus incorrect

(smooths are centered, so the factorial predictor also needs to be included in the fixed effects)

```
> m2 = bam(uV ~ s(Time, by=Correctness) + Correctness +
+             s(Time, Subject, bs='fs', m=1),
+             data=dat, gc.level=2, method='ML')
> summary(m2)
```

...

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.55608	0.29553	-1.882	0.0599 .
Correctnessincor	0.88609	0.05493	16.132	<2e-16 ***

Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
s(Time):Correctnesscor	1.295	1.458	2.901	0.0704 .
s(Time):Correctnessincor	4.220	5.145	26.494	<2e-16 ***
s(Time, Subject)	229.770	602.000	14.490	<2e-16 *****

R-sq. (adj) = 0.0205 Deviance explained = 2.1%
-ML = 1.9119e+06 Scale est. = 333.26 n = 442160

- ▶ Does the correct-incorrect distinction improve the model?

Assessing correct versus incorrect

(smooths are centered, so the factorial predictor also needs to be included in the fixed effects)

```
> m2 = bam(uV ~ s(Time, by=Correctness) + Correctness +
           s(Time, Subject, bs='fs', m=1),
           data=dat, gc.level=2, method='ML')
> summary(m2)

...
Parametric coefficients:
                        Estimate Std. Error t value Pr(>|t|)
(Intercept)      -0.55608    0.29553 -1.882   0.0599 .
Correctnessincor  0.88609    0.05493 16.132 <2e-16 ***
Approximate significance of smooth terms:
                    edf Ref.df     F p-value
s(Time):Correctnesscor  1.295  1.458  2.901  0.0704 .
s(Time):Correctnessincor  4.220  5.145 26.494 <2e-16 ***
s(Time, Subject)        229.770 602.000 14.490 <2e-16 ****
R-sq. (adj) =  0.0205  Deviance explained =  2.1%
-ML = 1.9119e+06  Scale est. = 333.26   n = 442160
```

- ▶ Does the correct-incorrect distinction improve the model?

Model comparison and GAMs

- ▶ There are three possibilities for comparing GAM models:
 - ▶ Compare (fRE)ML values: `compareML(m0, m1)` (custom function)
 - ▶ Compare AIC values (at least reduction of 2): `compareML(m0, m1)` or $AIC(m0) - AIC(m1)$
 - ▶ Look at the *p*-values of the model summary (**recommended**)
- ▶ **Important:** not all methods can be applied in all situations!
- ▶ The table below shows when each method can be applied:
(N.B. rho is used to correct for autocorrelation in the residuals and explained later)

	Gaussian without rho	Gaussian with rho	Non-Gaussian
gam	(RE)ML / AIC / summary	-	(RE)ML / AIC / summary
bam	(fRE)ML / AIC / summary	(fRE)ML / summary	AIC / summary

Assessing model improvement

```
> source('compareML.R')    # custom model comparison function
> compareML(m1,m2)

m1: uV ~ s(Time) + s(Time, Subject, bs = "fs", m = 1)

m2: uV ~ s(Time, by = Correctness) + Correctness +
     s(Time, Subject, bs = "fs", m = 1)

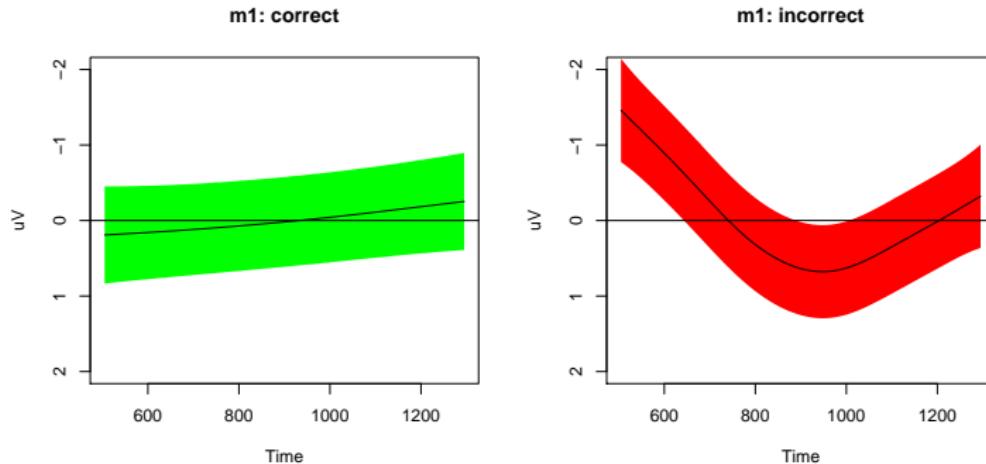
Chi-square test of ML scores
-----
   Model   Score  Edf   Chisq      Df  p.value Sig.
1     m1 1912123    5
2     m2 1911925    8 197.563 3.000 < 2e-16 ***

AIC difference: 393.34, model m2 has lower AIC.
```

- ▶ Model m2 is much better (AIC decrease > 2)!

Visualizing both patterns: partial effects

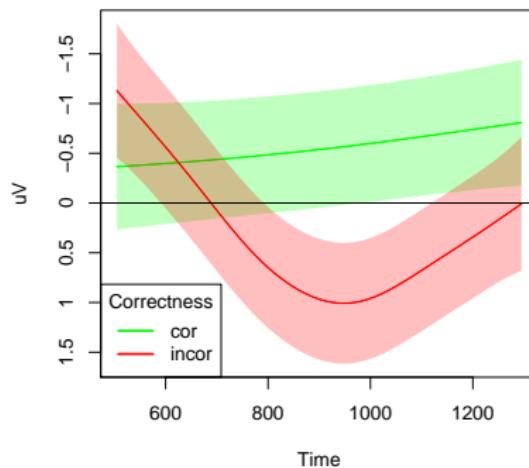
```
> par(mfrow=c(1,2))
> plot(m2, select=1, shade=T, rug=F, ylim=c(2,-2), shade.col='green',
       seWithMean=T, ylab='uv', main='m1: correct'); abline(h = 0)
> plot(m2, select=2, shade=T, rug=F, ylim=c(2,-2), shade.col='red',
       seWithMean=T, ylab='uv', main='m1: incorrect'); abline(h = 0)
```



Visualizing both patterns: complete model context

```
> plotSmooths(m2, "Time", "Correctness", colors=c('green','red'),  
              dropRanef='Subject')  
# Note that if there are additional variables in the model, plotSmooths  
# sets these to the specified value (via the 'cond' parameter) or to their  
# mean / most frequent value. Random-effect factors are excluded via  
# dropRanef.
```

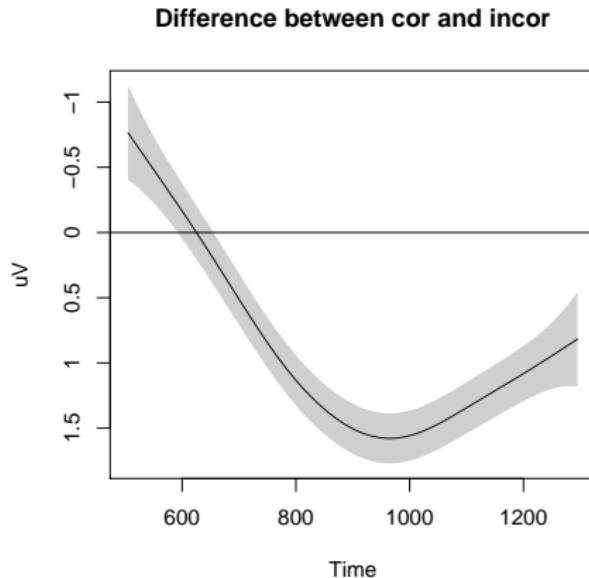
Model: m2



Visualizing the differences

(no individual variation w.r.t. the difference is taken into account: confidence bands **too** thin)

```
> plotDiff(m2, "Time", "Correctness") # custom function
```



Including individual variation for correct vs. incorrect

```
> dat$SubjectCor = interaction(dat$Subject,dat$Correctness) # new factor
> m3 = bam(uV ~ s(Time,by=Correctness) + Correctness +
+           s(Time, SubjectCor,bs='fs',m=1),
+           data=dat, gc.level=2, method='ML')
> summary(m3)

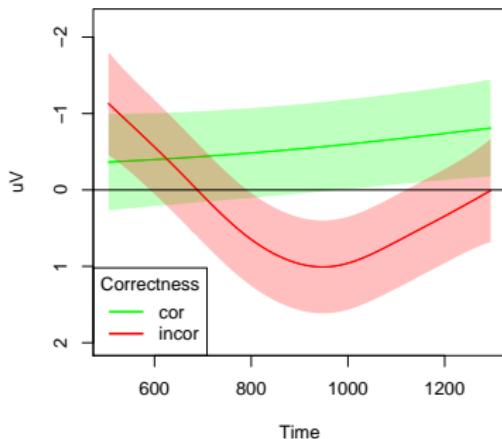
...
Parametric coefficients:
                        Estimate Std. Error t value Pr(>|t|)
(Intercept)          -0.5538     0.3974  -1.393   0.164
Correctnessincor    0.8965     0.5638   1.590   0.112

Approximate significance of smooth terms:
                               edf Ref.df      F p-value
s(Time):Correctnesscor    1.316  1.449  0.965  0.339
s(Time):Correctnessincor  4.125  4.939 17.724 <2e-16 ***
s(Time, SubjectCor)      469.508 1204.000 13.266 <2e-16 ***
R-sq.(adj) =  0.036  Deviance explained =  3.7% # was 2.1%
-ML = 1.9087e+06  Scale est. = 327.98    n = 442160
```

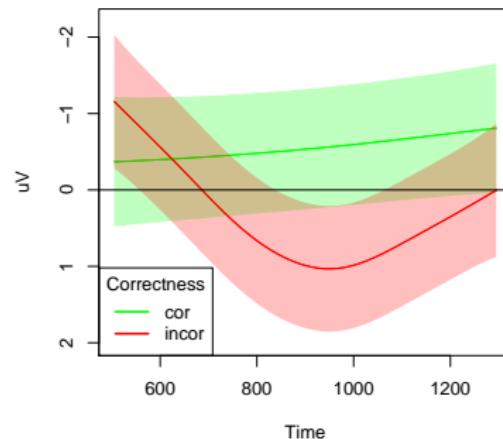
More uncertainty in the patterns

```
> par(mfrow=c(1,2))
> plotSmooths(m2, "Time", "Correctness", colors=c('green','red'),
  dropRanef='Subject', ylim=c(2,-2.2))
> plotSmooths(m3, "Time", "Correctness", colors=c('green','red'),
  dropRanef='SubjectCor', ylim=c(2,-2.2))
```

Model: m2

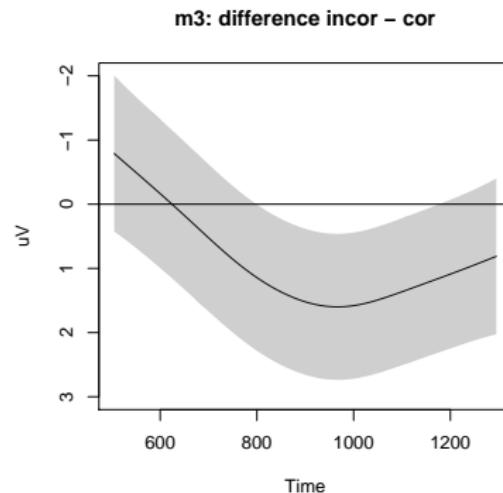
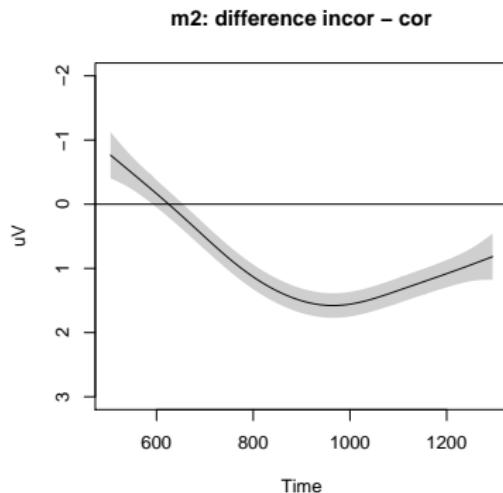


Model: m3



More uncertainty in the difference curve

```
> par(mfrow=c(1,2))
> plotDiff(m2, "Time", "Correctness", ylim=c(3,-2),
           main="m2: difference incor - cor")
> plotDiff(m3, "Time", "Correctness", ylim=c(3,-2),
           main="m3: difference incor - cor")
```



Explicit test for significant difference between levels

(smooths for a binary by-variable are **not** centered, so no extra term in the fixed effects)

```
# Note that the reference level influences the model. It makes sense
# to choose the reference level on the basis of your research question.
> dat$IsIncorrect = (dat$Correctness == 'incor')*1
> m3b = bam(uV ~ s(Time) + s(Time,by=IsIncorrect) +
+             s(Time,SubjectCor,bs='fs',m=1),
+             data=dat, gc.level=2, method='ML')
> summary(m3b)

...
Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.5486     0.3983   -1.377    0.168

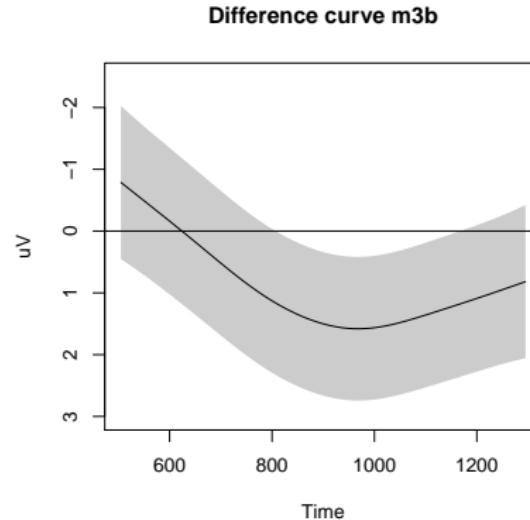
Approximate significance of smooth terms:
            edf Ref.df      F p-value
s(Time)       1.353  1.495  0.913  0.358
s(Time):IsIncorrect 4.947  5.724 12.007 9.27e-13 ***
s(Time,SubjectCor) 469.375 1204.000 13.266 < 2e-16 ***

R-sq.(adj) =  0.036  Deviance explained =  3.7%
-ML = 1.9087e+06  Scale est. = 327.98    n = 442160
```

Visualizing correct and incorrect difference

(identical shape as the difference curve of model m3)

```
> plot(m3b, shade=T, rug=F, ylim=c(1.5,-3.5), select=2,  
       main='Difference curve m3b', ylab='uV', seWithMean=T)  
> abline(h = 0)
```



Separating intercept and smooth difference

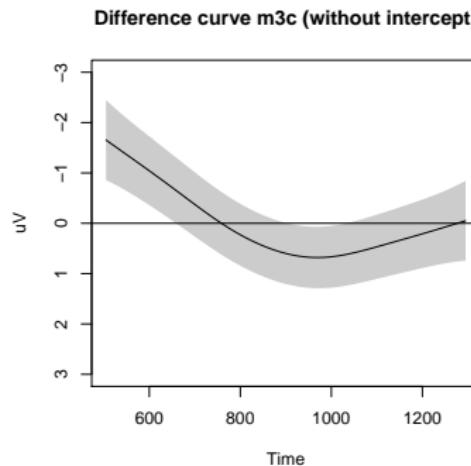
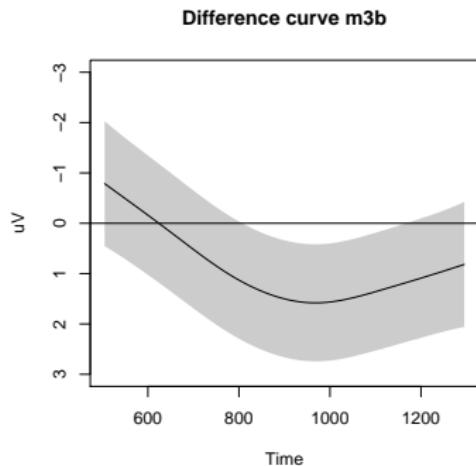
(the binary by-variable combines the intercept and smooth difference)

```
> dat$Correctness0 = as.ordered(dat$Correctness)
> contrasts(dat$Correctness0) = 'contr.treatment' # contrast treatment
> m3c = bam(uV ~ s(Time) + s(Time,by=Correctness0) + Correctness0 +
+             s(Time,SubjectCor,bs='fs',m=1),
+             data=dat, gc.level=2, method='ML')
> summary(m3c)
...
Parametric coefficients:
                        Estimate Std. Error t value Pr(>|t|)
(Intercept)           -0.5380    0.3979  -1.352   0.176
Correctness0incor     0.8799    0.5640   1.560   0.119

Approximate significance of smooth terms:
                     edf Ref.df      F p-value
s(Time)              1.572 1.783  1.02  0.342
s(Time):Correctness0incor 3.968 4.748 12.80 1.1e-11 ***
s(Time,SubjectCor)   469.380 1204.000 13.27 < 2e-16 ***
R-sq.(adj) = 0.036  Deviance explained = 3.7%
-ML = 1.9087e+06  Scale est. = 327.98    n = 442160
```

Visualizing difference between correct and incorrect

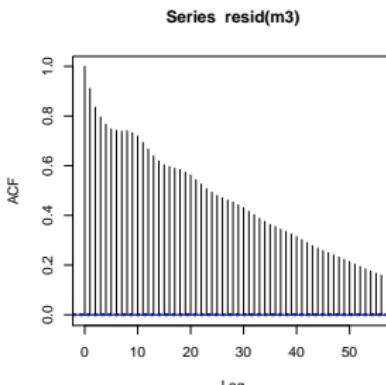
```
> plot(m3b, shade=T, rug=F, ylim=c(3,-4), select=2, ylab='uV',  
       main='Difference curve m3b', seWithMean=T)  
> abline(h = 0)  
> plot(m3c, shade=T, rug=F, ylim=c(3,-4), select=2, ylab='uV',  
       main='Difference curve m3c (without intercept)', seWithMean=T)  
> abline(h = 0)
```



Autocorrelation in the data is a huge problem!

(residuals should be independent, otherwise the standard errors and p -values are wrong)

```
# It is essential that the data used in m3 is sorted per individual EEG
# signal sequence (per subject, per trial, per ROI, etc.). The best rho
# value can be found by trying different values and comparing models. (The
# autocorrelation at lag 1 for the same model w/o rho seems to work well,
# however.) A good approach is to use the lowest value possible which has
# a good acf plot and is not significantly worse than the best rho model.
> m3acf = acf(resid(m3)) # show autocorrelation
> rhoval = as.vector(m3acf[1]$acf)
> rhoval
[1] 0.91104 # correlation of residuals at time t with those at time t-1
```



Correcting for autocorrelation

(rho can only be used with bam, not with gam)

```
> m4 = bam(uV ~ s(Time,by=Correctness) + Correctness +
           s(Time,SubjectCor,bs='fs',m=1), data=dat,
           gc.level=2, method='ML', rho=rhoval, AR.start=SeqStart)
> summary(m4)

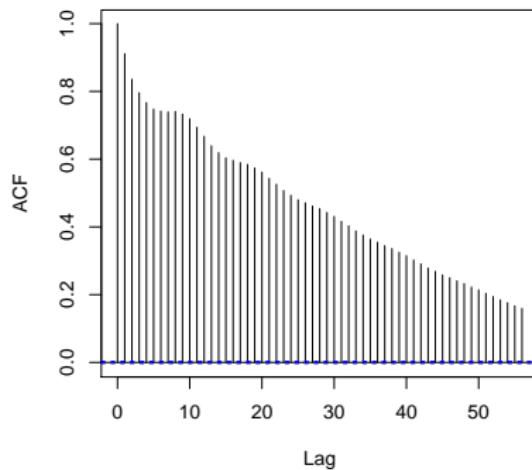
...
Parametric coefficients:
                        Estimate Std. Error t value Pr(>|t|)
(Intercept)      -0.5453     0.4054  -1.345   0.179
Correctnessincor  0.8932     0.5746   1.554   0.120

Approximate significance of smooth terms:
                               edf Ref.df      F p-value
s(Time):Correctnesscor    1.011  1.022 0.323  0.575
s(Time):Correctnessincor  3.099  4.028 6.859 1.56e-05 ***
s(Time,SubjectCor)       116.095 1204.000 0.657 < 2e-16 ***
R-sq.(adj) = 0.0321 Deviance explained = 3.24%
-ML = 1.4923e+06 Scale est. = 287.48 n = 442160
```

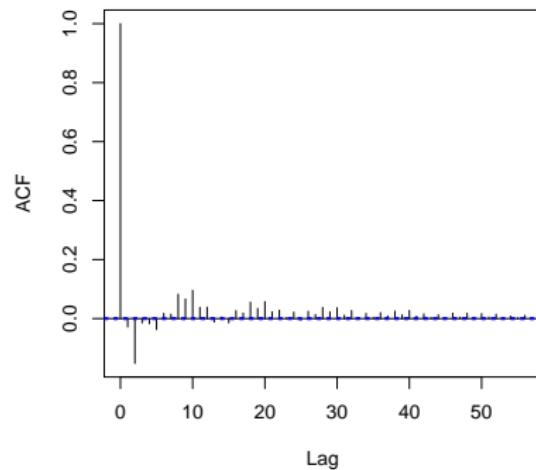
Autocorrelation has been removed

```
> acf.new(m3, m4, "Time") # custom plotting function
```

Original acf



New acf with rho: 0.91



Clear model improvement

```
> compareML(m3,m4)

m3: uV ~ s(Time, by = Correctness) + Correctness +
     s(Time, SubjectCor, bs = "fs", m = 1)

m4: uV ~ s(Time, by = Correctness) + Correctness +
     s(Time, SubjectCor, bs = "fs", m = 1)

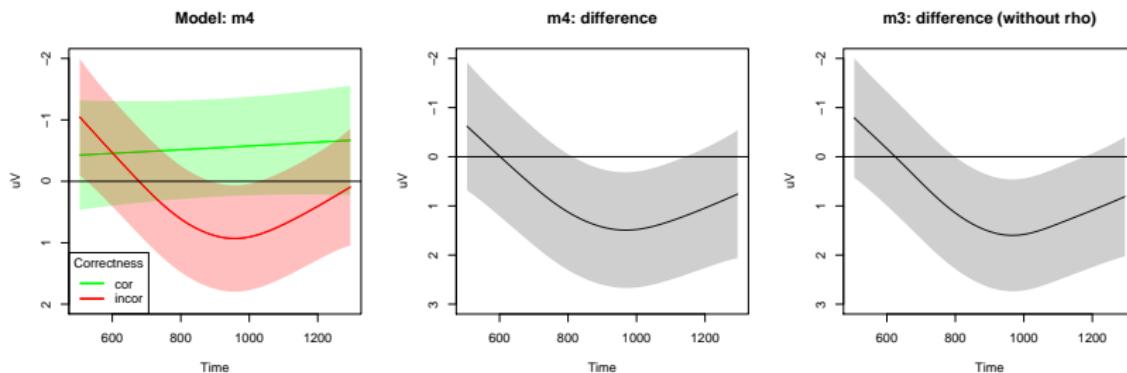
Chi-square test of ML scores
-----
   Model   Score Edf      Chisq    Df p.value Sig.
1     m3 1908692    8 416431.071 0.000 < 2e-16 ***
2     m4 1492261    8

Warning message:
In compareML(m3, m4) :
  AIC is not reliable, because an AR1 model is included ...

> AIC(m3) - AIC(m4) # AIC indeed does not work when using rho
[1] -1428.872
```

Resulting smooths: slightly wider confidence bands

```
> par(mfrow=c(1,3))
> plotSmooths(m4, "Time", "Correctness", colors=c('green','red'),
+ dropRanef="SubjectCor", ylim=c(2,-2))
> plotDiff(m4, "Time", "Correctness", ylim=c(3,-2), main='m4: difference')
> plotDiff(m3, "Time", "Correctness", ylim=c(3,-2),
+ main='m3: difference (without rho)')
```



Our research question: the effect of age of arrival

```
# in this lecture, rho is held constant at 0.91
> m5 = bam(uV ~ te(Time,AoArr,by=Correctness) + Correctness +
+             s(Time,SubjectCor,bs='fs',m=1), data=dat,
+             gc.level=2, method='ML', rho=rhoval, AR.start=SeqStart)

> summary(m5)

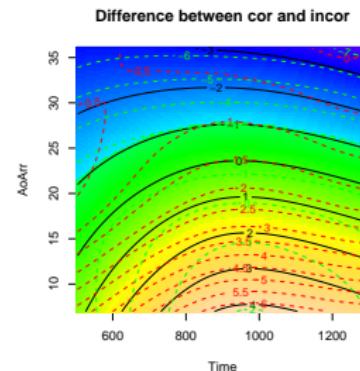
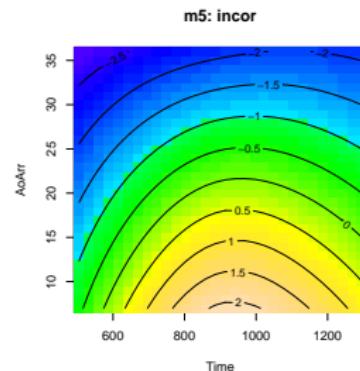
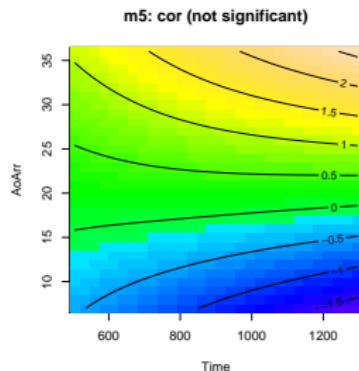
...
Parametric coefficients:
                               Estimate Std. Error t value Pr(>|t|)
(Intercept)           -0.5589     0.3962  -1.411   0.1584
Correctnessincor     0.9260     0.5615   1.649   0.0991 .

Approximate significance of smooth terms:
                                         edf Ref.df F p-value
te(Time,AoArr):Correctnesscor      3.250 3.482 1.565 0.186
te(Time,AoArr):Correctnessincor  5.871 6.956 4.707 2.93e-05 ***
s(Time,SubjectCor)        113.435 1202.000 0.628 < 2e-16 ***
R-sq.(adj) = 0.0322 Deviance explained = 3.25%
-ML = 1.4923e+06 Scale est. = 287.48 n = 442160
```

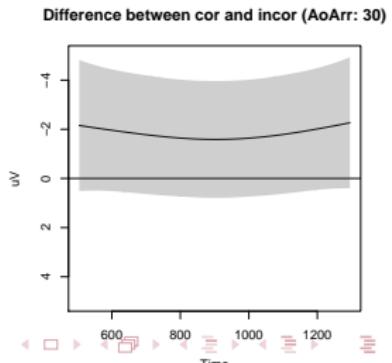
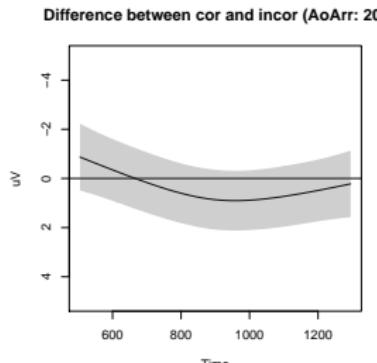
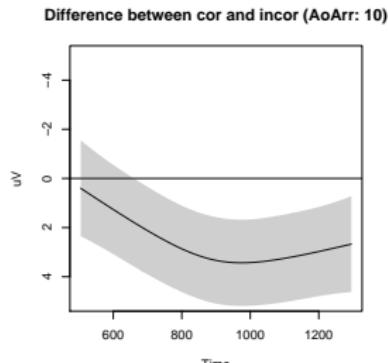
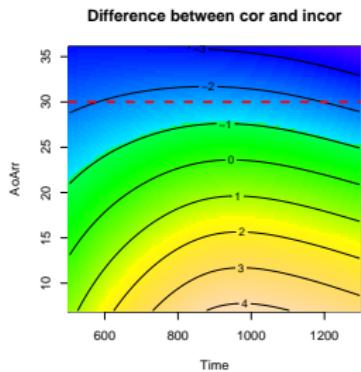
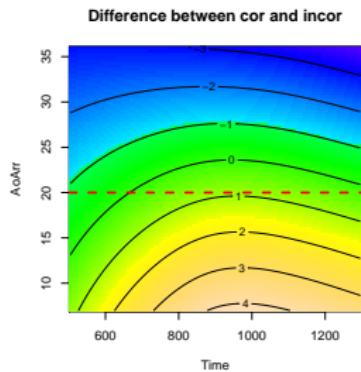
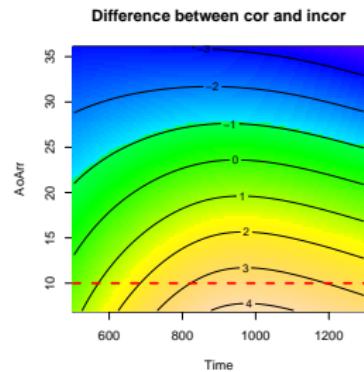
Visualization of the 2-dimensional smooths

```
# Custom plotting functions are used to plot the partial effects
# (i.e. the general pattern) and difference surfaces.
# N.B. a surface including all effects (at pre-specified levels) can be
# visualized using vis.gam(...) - see additional slides at the end.

> par(mfrow=c(1,3))
> pvis.gam(m5, plot.type='contour', view=c('Time','AoArr'), select=1,
           color='topo', main='m5: cor (not significant)')
> pvis.gam(m5, plot.type='contour', view=c('Time','AoArr'), select=2,
           color='topo', main='m5: incor')
> plotDiff2D(m5, "Time", "AoArr", "Correctness", plotCI=T)
```



Interpreting 2-dimensional smooths



Significance testing using a binary variable

(an ordered factor may be used to separate the intercept difference from the difference surface)

```
> m5b = bam(uV ~ te(Time,AoArr) + te(Time,AoArr,by=IsIncorrect) +
+             s(Time,SubjectCor,bs='fs',m=1), data=dat,
+             gc.level=2, method='ML', rho=rhoval, AR.start=SeqStart)
> summary(m5b)
...
Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.5222     0.4004   -1.304    0.192

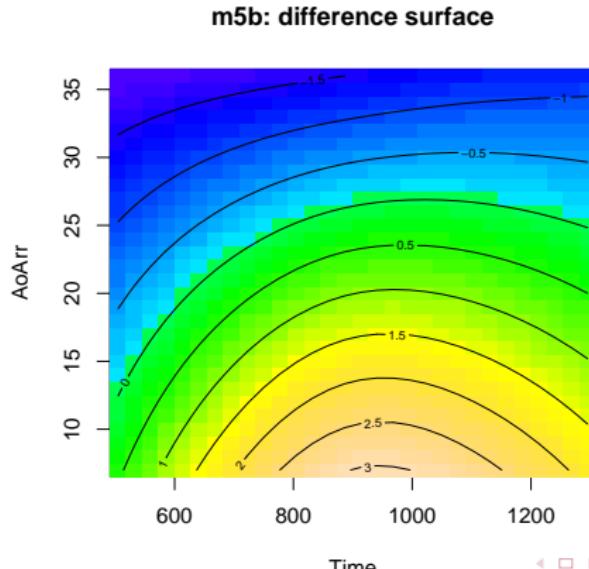
Approximate significance of smooth terms:
                               edf Ref.df      F p-value
te(Time,AoArr)           1.088 20.000 0.083 0.12947
te(Time,AoArr):IsIncorrect 6.610  7.631 3.402 0.00087 ***
s(Time,SubjectCor)       114.647 1203.000 0.636 < 2e-16 ***

R-sq.(adj) =  0.0321  Deviance explained = 3.24%
-ML = 1.4923e+06  Scale est. = 287.48    n = 442160
```

Visualizing the difference surface

(note that m5b is a different model than m5, so the difference surfaces do not **need** to be identical)

```
> pvis.gam(m5b, plot.type='contour', view=c('Time','AoArr'),  
           contour.col='black', select=2, color='topo',  
           main='m5b: difference surface')
```



Assessing if k needs to be increased

(note that a single k value specifies the wigglyness per dimension in te() / in total in s())

```
# the wigglyness may also specified separately per dimension for te()
> m5c = bam(uV ~ te(Time,AoArr,by=Correctness,k=c(10,10)) + Correctness +
+             s(Time,SubjectCor,bs='fs',m=1), data=dat,
+             gc.level=2, method='ML', rho=rhoval, AR.start=SeqStart)
> compareML(m5,m5c) # the additional complexity is not needed
...
Chi-square test of ML scores
-----
Model      Score   Edf  Chisq     Df  p.value Sig.
1    m5c  1492256   14
2      m5  1492255   14  0.271  0.000 < 2e-16 ***

...
> summary(m5c) # max edf: 99 (10^2 - 1)
...
edf  Ref.df      F  p-value
te(Time,AoArr):Correctnesscor  3.239  3.465 1.644    0.167    # was 3.3
te(Time,AoArr):Correctnessincor 6.397  8.010 4.072 7.31e-05 *** # was 5.9
...
```

Adding factor smooths per word(-correctness)

```
> dat$WordCor = interaction(dat$Word,dat$Correctness)
> m6 = bam(uV ~ te(Time,AoArr,by=Correctness) +
+            Correctness + s(Time,SubjectCor,bs='fs',m=1) +
+            s(Time,WordCor,bs='fs',m=1), data=dat, gc.level=2,
+            method='ML', rho=rhoval, AR.start=SeqStart) # 15 hours...
> summary(m6)
...
Parametric coefficients:
                               Estimate Std. Error t value Pr(>|t|)
(Intercept)           -0.5888     0.5571  -1.057   0.291
Correctnessincor     0.7473     0.7804   0.958   0.338

Approximate significance of smooth terms:
                                         edf Ref.df      F p-value
te(Time,AoArr):Correctnesscor    3.106  3.207 1.774  0.145
te(Time,AoArr):Correctnessincor 5.864  6.943 4.793 2.32e-05 ***
s(Time,SubjectCor)             111.028 1202.000 0.624 < 2e-16 ***
s(Time,WordCor)                133.202 1726.000 0.260 < 2e-16 ***
R-sq.(adj) =  0.0511  Deviance explained = 5.17% # was 3.25%
-ML = 1.4921e+06  Scale est. = 287.19    n = 442160
```

Model comparison for factor smooths

```
> compareML(m5r, m6r)

# m5r and m6r are the same as m5 and m6, but with method='REML'
# (required when comparing random effects)

m5r: uV ~ te(Time, AoArr, by = Correctness) + Correctness +
      s(Time, SubjectCor, bs = "fs", m = 1)
m6r: uV ~ te(Time, AoArr, by = Correctness) + Correctness +
      s(Time, SubjectCor, bs = "fs", m = 1) +
      s(Time, WordCor, bs = "fs", m = 1)

Chi-square test of REML scores
-----
  Model   Score   Edf    Chisq     Df p.value Sig.
1  m5r  1492248    14
2  m6r  1492137    16 110.657  2.000 < 2e-16 ***
```

Decomposition: the pure effect of age of arrival

(note that the difference between the 2 ti-surfaces is also non-significant, not shown)

```
> m6b = bam(uV ~ s(Time,by=Correctness) + s(AoArr,by=Correctness) +
+   ti(Time,AoArr,by=Correctness) + Correctness +
+   s(Time,SubjectCor,bs='fs',m=1) +
+   s(Time,WordCor,bs='fs',m=1), data=dat, gc.level=2,
+   method='ML', rho=rhoval, AR.start=SeqStart)
> summary(m6b)
...

```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.5790	0.4606	-1.257	0.209
Correctnessincor	0.8279	0.6527	1.269	0.205

Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
s(Time):Correctnesscor	1.003	1.006	0.322	0.5715
s(Time):Correctnessincor	3.100	4.029	6.866	1.54e-05 ***
s(AoArr):Correctnesscor	1.002	1.002	2.883	0.0894 .
s(AoArr):Correctnessincor	1.002	1.002	4.268	0.0388 *
ti(Time,AoArr):Correctnesscor	1.263	1.491	2.059	0.1346
ti(Time,AoArr):Correctnessincor	1.390	1.723	0.163	0.8177
s(Time,SubjectCor)	111.134	1202.000	0.624	< 2e-16 ***
s(Time,WordCor)	133.225	1726.000	0.260	< 2e-16 ***

A simpler model without the non-linear interaction

```
> m6c = bam(uV ~ s(Time,by=Correctness) + s(AoArr,by=Correctness) +
+             Correctness + s(Time,SubjectCor,bs='fs',m=1) +
+             s(Time,WordCor,bs='fs',m=1), data=dat, gc.level=2,
+             method='ML', rho=rhoval, AR.start=SeqStart)
> summary(m6c)
...
Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.5789    0.4607 -1.257   0.209
Correctnessincor  0.8279    0.6527  1.268   0.205

Approximate significance of smooth terms:
          edf Ref.df     F p-value
s(Time):Correctnesscor 1.008  1.016 0.329  0.5700
s(Time):Correctnessincor 3.101  4.030 6.868 1.53e-05 ***
s(AoArr):Correctnesscor 1.009  1.010 2.793  0.0941 .
s(AoArr):Correctnessincor 1.013  1.015 4.131  0.0415 *
s(Time,SubjectCor)      111.255 1202.000 0.623 < 2e-16 ***
s(Time,WordCor)         134.127 1726.000 0.260 < 2e-16 ***

R-sq.(adj) =  0.0511  Deviance explained = 5.16%
-ML = 1.4921e+06  Scale est. = 287.19    n = 442160
```

Model comparison: the simpler model is sufficient

```
> compareML(m6b,m6c)
```

```
m6b: uV ~ s(Time, by = Correctness) + s(AoArr, by = Correctness) +
      ti(Time, AoArr, by = Correctness) + Correctness +
      s(Time, SubjectCor, bs = "fs", m = 1) +
      s(Time, WordCor, bs = "fs", m = 1)
```

```
m6c: uV ~ s(Time, by = Correctness) + s(AoArr, by = Correctness) +
      Correctness + s(Time, SubjectCor, bs = "fs", m = 1) +
      s(Time, WordCor, bs = "fs", m = 1)
```

Chi-square test of ML scores

```
-----
  Model   Score Edf Chisq    Df p.value Sig.
1  m6c  1492148   14
2  m6b  1492146   20 1.568  6.000    0.792
```

Warning messages:

```
1: In compareML(m6b, m6c) : AIC is not reliable, ...
2: In compareML(m6b, m6c) : Only small difference in ML...
```

Testing for significant smooth differences

(a binary variable may only occur once in a model, so ordered factors are essential here)

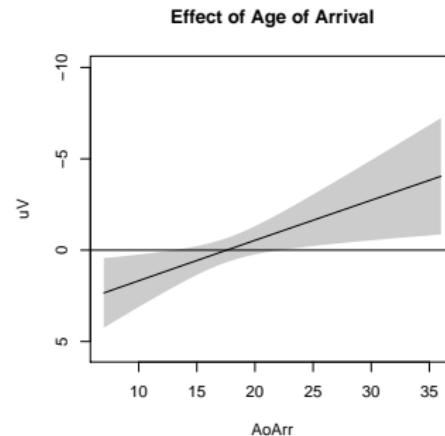
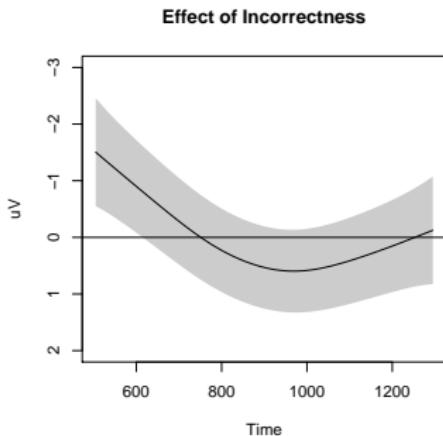
```
> m7 = bam(uV ~ s(Time) + s(Time,by=CorrectnessO) + s(AoArr) +
+           s(AoArr,by=CorrectnessO) + CorrectnessO + ...)
> summary(m7)
...
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.5768     0.4607 -1.252   0.211
CorrectnessOincor  0.8262     0.6527  1.266   0.206

Approximate significance of smooth terms:
                      edf Ref.df F p-value
s(Time)                 1.055 1.102 0.389 0.55289
s(Time):CorrectnessOincor 3.096 4.023 6.118 6.26e-05 ***
s(AoArr)                1.000 1.000 2.801 0.09423 .
s(AoArr):CorrectnessOincor 1.013 1.015 6.837 0.00871 **
s(Time,SubjectCor)       110.670 1202.000 0.623 < 2e-16 ***
s(Time,WordCor)          134.023 1726.000 0.260 < 2e-16 ***
...
```

- ▶ Thus: the data clearly show a non-linear effect of time and age of arrival

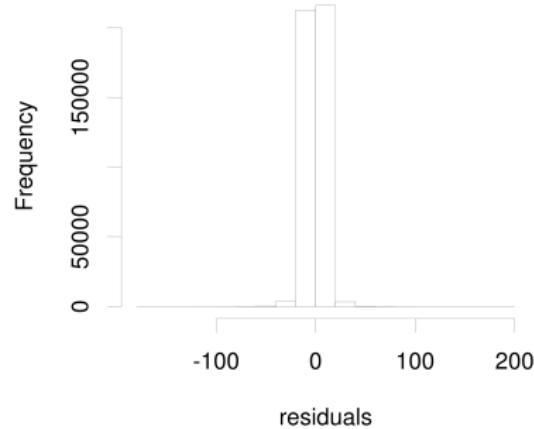
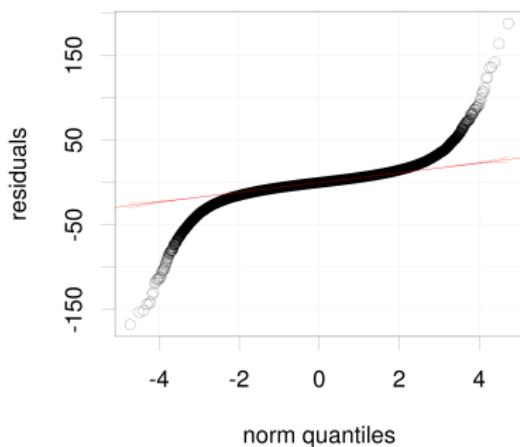
Difference curves (smooths only)

```
> par(mfrow=c(1,2))
> plot(m7, shade=T, rug=F, select=2, ylim=c(2,-3),
       main='Effect of Incorrectness', ylab='uV', seWithMean=T)
> abline(h = 0)
> plot(m7, shade=T, rug=F, select=4, ylim=c(5.5,-10),
       main='Effect of Age of Arrival', ylab='uV', seWithMean=T)
> abline(h = 0)
```



Finally: model criticism

```
> library(car)
> par(mfrow=c(1,2))
> qqplot.rho(m7) # custom qqplot of residuals taking rho into account
> hist.rho(m7) # custom histogram of residuals taking rho into account
```



Approach to normalize residuals

- ▶ This type of residual distribution is common for EEG data
- ▶ These deviations are **very problematic**: incorrect p -values...
- ▶ Solution: model fitting using scaled- t distribution
 - ▶ Implemented in `gam`: `family = "scat"`
 - ▶ Adding `rho` to `gam` possible via custom function (created by Natalya Pya)
 - ▶ Unfortunately, `gam` is too slow, so a `bam` approach is in development
 - ▶ Current state: implemented for `fREML` estimation with pre-specified scaled- t parameters (θ)
 - ▶ Thus `gam` still necessary to determine θ (but applied to subset / simpler model specification)

Estimating scaled-t parameters efficiently

1. Fit an intercept-only scaled-*t* gam model (including rho, method="REML") to all data and extract θ
2. Fit the complete model without random effects as a scaled-*t* gam model (including rho, method="REML") to all data and extract θ
3. Compare the models' θ values:
 - ▶ Relatively small differences: use θ of the complex model in the bam estimation
 - ▶ Relatively large differences: determine θ by fitting the complete model specification (including random effects) to a small subset of the data (randomly select trials, compare θ across a few runs)

Fitting a scaled-*t* model

(does not work for binomial models)

```
> source('bam.art.fit.R') # custom functions by Natalya Pya
> g0 <- gam(uV ~ 1, data=dat, method='REML',
+             family=scat(rho=rhoval, AR.start=dat$SeqStart))
> g1 <- gam(uV ~ s(Time) + s(Time,by=Correctness0) + s(AoArr) +
+             s(AoArr,by=Correctness0) + Correctness0,
+             data=dat, method='REML',
+             family=scat(rho=rhoval, AR.start=dat$SeqStart)))
> theta0 = g0$family$getTheta(TRUE)
> thetal = g1$family$getTheta(TRUE) # 4.566764 5.264096
> theta0 - thetal
[1] -0.003382087 -0.002104720 # OK, not too high

# for simplicity and speed we ignore by-word factor smooths here
# (those should be included in the final model, however)
> m8 = bam(uV ~ s(Time) + s(Time,by=Correctness0) + s(AoArr) +
+             s(AoArr,by=Correctness0) + Correctness0 +
+             s(Time,SubjectCor,bs='fs',m=1), data=dat, gc.level=2,
+             method='fREML', family=art(theta=thetal,rho=rhoval),
+             AR.start=SeqStart)) # 12 hours (35 min. for Gaussian)!

# for a fair comparison a Gaussian model (m8g) was fitted using fREML
```

Using the scaled-*t* distribution: *p*-values change

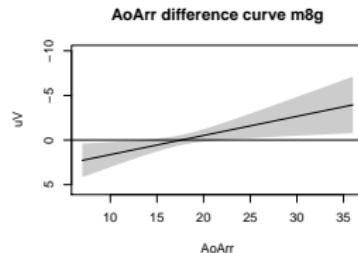
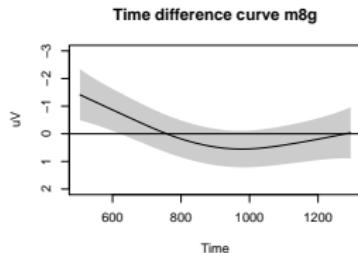
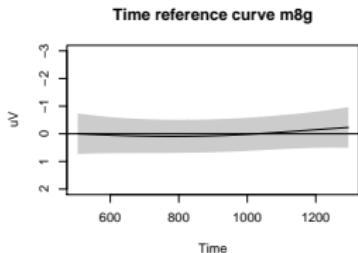
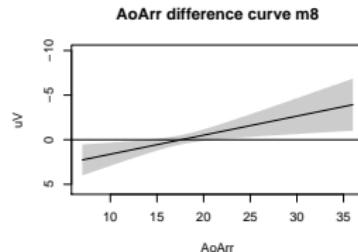
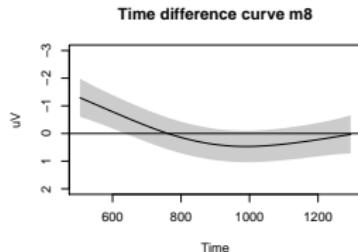
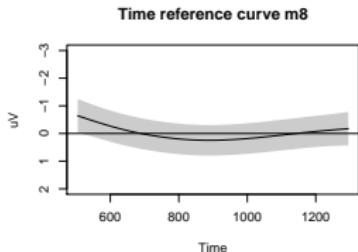
```
> summary(m8)
```

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-0.002761	0.369947	-0.007	0.994	
CorrectnessOincor	0.842288	0.523107	1.610	0.107	
s(Time)	3.600	4.643	4.280	0.00107	**
s(Time):CorrectnessOincor	3.475	4.485	7.940	8.71e-07	***
s(AoArr)	1.692	1.714	1.164	0.29437	
s(AoArr):CorrectnessOincor	1.001	1.001	7.496	0.00617	**
s(Time, SubjectCor)	122.258	1202.000	1.795	< 2e-16	***
R-sq. (adj) = 0.0291	Deviance explained = 1.9%				

```
> summary(m8g)
```

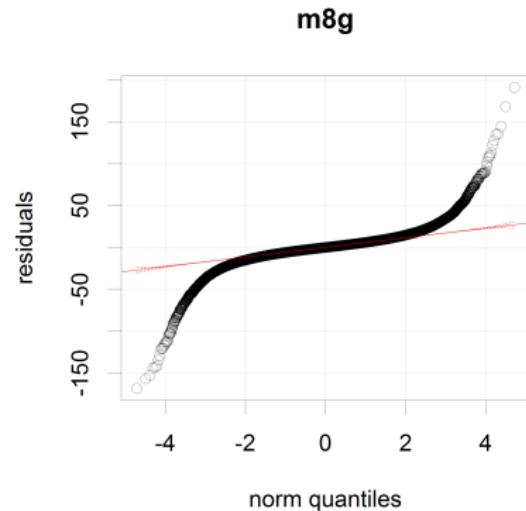
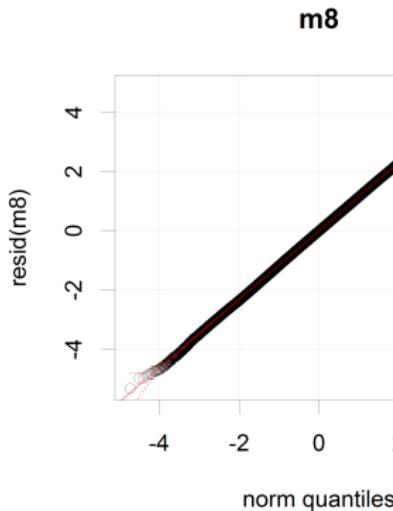
	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-0.5370	0.4006	-1.34	0.18	
CorrectnessOincor	0.9076	0.5672	1.60	0.11	
s(Time)	1.660	2.085	0.577	0.56858	
s(Time):CorrectnessOincor	3.058	3.968	4.414	0.00152	**
s(AoArr)	1.698	1.758	1.501	0.21636	
s(AoArr):CorrectnessOincor	1.004	1.004	6.431	0.01114	*
s(Time, SubjectCor)	110.461	1202.000	0.609	< 2e-16	***
R-sq. (adj) = 0.0321	Deviance explained = 3.24%				

Using the scaled- t distribution: patterns remain similar



Model criticism: much improved!

```
> library(car)
> par(mfrow=c(1,2))
> qqp(resid(m8), main='m8') # resid of scaled-t model takes rho into account
> qqplot.rho(m8g, main='m8g') # custom qqplot of resid. accounting for rho
```



Discussion

- ▶ Still much to do: e.g., testing the significance of other possibly important variables (proficiency, age, etc.)
- ▶ But don't make it too complex!
 - ▶ There is much variation present in EEG data and adding very complex surfaces will almost certainly improve your model significantly
 - ▶ keep it simple, otherwise you won't be able to compute/interpret the results
 - ▶ Also note that for hypothesis testing, it is essential not to conduct step-wise approaches

Conclusion

- ▶ GAMs are very useful to analyze EEG and other time-series data
 - ▶ The method is very suitable to detect **non-linear patterns**, while taking into account individual variation and correcting for autocorrelation
- ▶ If you would like to get some hands-on experience, you can use the lab session available here:

<http://www.let.rug.nl/wieling/statscourse/lecture4>

Thank you for your attention!

A red-footed booby stands on a dark, craggy rock. It has white plumage on its head, neck, and body, with black wing tips and a black patch around its eye. Its feet are a bright orange-red color. The background consists of dense green foliage and branches.

Any
questions?

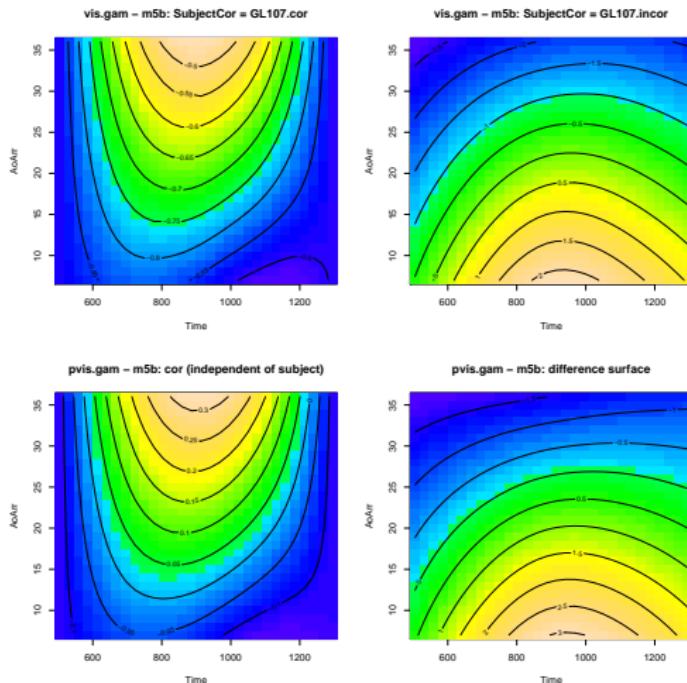
Visualization: vis.gam vs. pvis.gam (1)

```
> m5b = bam(uV ~ te(Time,AoArr) + te(Time,AoArr,by=IsIncorrect) +
+             s(Time,SubjectCor,bs='fs',m=1), data=dat,
+             gc.level=2, method='ML', rho=rhoval, AR.start=SeqStart)

> par(mfrow=c(2,2))
> vis.gam(m5b, plot.type='contour', view=c('Time','AoArr'),
+           cond=list(IsIncorrect=0,Subject='GL107.cor'),
+           contour.col='black', color='topo',
+           main='vis.gam - m5b: SubjectCor = GL107.cor')
> vis.gam(m5b, plot.type='contour', view=c('Time','AoArr'),
+           cond=list(IsIncorrect=1,Subject='GL107.incor'),
+           contour.col='black', color='topo',
+           main='vis.gam - m5b: SubjectCor = GL107.incor')
> pvis.gam(m5b, plot.type='contour', view=c('Time','AoArr'),
+            contour.col='black', select=1, color='topo',
+            main='pvis.gam - m5b: cor (independent of subject)')
> pvis.gam(m5b, plot.type='contour', view=c('Time','AoArr'),
+            contour.col='black', select=2, color='topo',
+            main='pvis.gam - m5b: difference surface')
```

Visualization: vis.gam vs. pvis.gam (2)

(vis.gam shows the resulting surface of the conditions, pvis.gam shows the smooths in your model)



Differences between ML, REML and GCV

- ▶ ML is conservative, but the variance component (i.e. smooth) is biased (oversmoothed)
- ▶ (f) REML is less conservative, but robust to moderate autocorrelation, and the smooth is not biased
- ▶ GCV.Cp is better for prediction, but not robust to autocorrelation
- ▶ Strongest approach is to compare results across different methods (i.e. using AIC/(RE)ML)

